

# Wasted Potential: The Value of a Life and the Significance of What Could Have Been\*

Michal Masny  
mmasny@princeton.edu  
26 October 2022

**Abstract:** According to the orthodox view, the goodness of a life depends exclusively on the things that actually happened within it, such as its pleasures and pains, the satisfaction of its subject's preferences, or the presence of various objective goods and bads. In this paper, I propose that the goodness of a life also depends on what could have happened, but didn't. I then argue that this view helps us resolve three ethical puzzles concerning the standards for a life worth living for animals, the significance of a life's shape, and the badness of death.

**Word count:** 9,800

## 1. Introduction

It is a striking fact about us that we care deeply about what could have happened, but didn't. This paper explores the significance of this concern for theorizing about the goodness of a life. To get a feel for the issues that we will be examining, consider the case of Sophie Germain, a French mathematician of the early 19<sup>th</sup> century. She was born to a wealthy Parisian family and enjoyed a life rich in meaningful relationships, sophisticated pleasures, and important achievements. However, much of her exceptional academic talent was wasted because of the obstacles she faced as a woman. Early on, her parents tried to hinder her youthful fascination with mathematics. Later, she was barred from attending the École Polytechnique and the meetings of the Paris Academy of Sciences, and both her manuscripts and published work were regularly ignored by her contemporaries.

Germain's case evokes an attitude of *evaluative ambivalence*. On the one hand, her story is uplifting: she experienced a lot of what makes life valuable. On the other, there is a sense of tragedy that we just cannot shrug off: she could have achieved much more, but didn't, and could have been more appreciated and spared many frustrations, but wasn't.

Situations which have this general structure uncover something important for value theory. According to the orthodox view, how good a life is for its subject depends exclusively on the things that actually happened within it, such as its pleasures and pains, the satisfaction of its subject's preferences, or the presence of various objective

---

\* *Draft*. Comments welcome. But please do not cite without permission.

goods and bads. In this paper, I challenge this orthodox view and argue that the goodness of a life also depends on what could have happened, but didn't.

My discussion has three parts. In the first part (§2-3), I present initial motivation for *the Dual Theory*, which holds that how good a life is for someone is determined jointly by their level of well-being and the degree to which they realize their potential. In the second part (§4-5), I examine different ways in which this theory can be spelled out. In particular, I consider what counts as an individual's potential and how well-being and the realization of one's potential combine to determine the overall goodness of a life. In the third part (§6-8), I provide additional support for the Dual Theory by highlighting its explanatory power. Specifically, I argue that it helps us resolve three ethical puzzles concerning the standards for a life worth living for non-human animals, the significance of a life's shape, and the badness of death.

Two preliminary remarks will make it easier for the reader to appreciate the scope of this discussion and its position in the broader theoretical landscape.

First, a few words about the key concepts. This paper is about the goodness of a life for its subject or, equivalently, the prudential value of a life. Many philosophers assume that there is only one dimension of the good life, well-being, which leads them to use these two terms as synonyms. For the present purposes, however, we want to use the 'the goodness of a life' as a more encompassing notion than 'well-being'. This terminological convention has an influential precedent: for example, Susan Wolf (2010), Antti Kauppinen (2012), and Aaron Smuts (2017) hold that how good a life is for a person is determined not just by their well-being, but also by how meaningful it is. This paper, too, contends that there are at least two dimensions along which we can assess the goodness of a life: well-being and the realization of one's potential.<sup>2</sup>

Second, some context. Although the idea sketched above has not yet been discussed in contemporary value theory, there are some thinkers who have come within a stone's throw of appreciating it. Jeff McMahan (1996) and Peter Vallentyne (2005), for instance, consider the idea that facts about an individual's potential for well-being bear on whether we have egalitarian reasons to redistribute limited resources to that individual. Moreover, Frick and Lerner (ms) suggest that similar considerations influence whether it is permissible to bring an individual into existence. Finally, R. Jay Wallace (2013) claims that what we have reasons to regret in our lives depends in part on what else could have happened.<sup>3</sup>

---

<sup>2</sup> I do not take a stance here on whether meaningfulness also affects the goodness of a life.

<sup>3</sup> See also Harman (2009) and Setiya (2016) for discussions of this issue.

If the arguments of this paper are sound, however, the significance of what could have been is more pervasive than these philosophers have contended. For one thing, this significance is not restricted to the first-personal, retrospective standpoint from which we tend to assess the aptness of attitudes like regret. For another, it is already present at the level of prudential value, and its bearing on deontic matters—such as those concerning the demands of justice and the reasons that apply to the creation of individuals—is at most derivative.

But let's start at the beginning, with a closer look at the phenomenon of evaluative ambivalence which we have identified earlier.

## 2. Evaluative ambivalence

Consider the case of Sophie Germain again. By the lights of any standard theory of well-being, she had an excellent life, even once we factor in the disappointments and the struggles. And yet it is clear that things could have gone much better for her. This makes her life seem at once heartening and tragic. Crucially, this evaluative ambivalence persists even as we continue to think long and hard about this case. Germain's life does not suddenly become just heartening or just tragic. A full and accurate description of the quality of her life seems to preclude merely summing the two evaluations, positive and negative. Instead, it requires that we maintain both judgments at once.

The enduring character of evaluative ambivalence in such cases leaves us with two options. The first is to resign ourselves to the idea that the judgments in question are inconsistent: a whole life cannot simultaneously be both very good and very bad. The second, more attractive option is to consider the possibility that these judgments are not in conflict but have different objects. This is the line of inquiry pursued in this paper.

The Dual Theory, I believe, gives an elegant explanation of our evaluative ambivalence in Germain's case. One of our judgments has as its object her well-being, whereas the other concerns the realization of her potential, and both of these are relevant to the overall goodness of her life.

To get a better grip on the strength of this explanation, let's consider some of its competitors. According to the first set of rival explanations, only one of the judgments comprising our evaluative ambivalence has to do with the prudential value of a life, whereas the other is about its *impersonal value*. There are two versions of this thought.

First, some thinkers maintain that certain things are valuable even if they are not good for anyone in particular: existence of beautiful artworks and landscapes, equality

between members of the moral community, or accumulation of knowledge. That is, in addition to prudential value, we should recognize aesthetic, egalitarian, or epistemic value.<sup>4</sup> Now, one might suggest that when we judge Germain's life to be in one respect tragic, we are making a judgment about one of these impersonal values. Perhaps a life which falls short of its potential makes for a worse story, and thus has low *aesthetic value*.

Second, problems of population ethics raised by Derek Parfit (1984) have led some authors to distinguish between the prudential value of a life and its *contributive value*. According to the critical level view, for example, a person's life contributes positively to the value of a state of affairs just in case its prudential value is sufficiently high, above a certain positive critical level.<sup>5</sup> Consequently, a life which is 'worth living' might turn out to be not 'worth creating'. In a similar vein, one might suggest that facts about what could have happened bear just on the contributive value of a life. Perhaps we find Germain's life to be in one respect tragic because, although her life is worth living in virtue of all the good things that happened to her, the fact that she fell short of her potential somehow makes her life not worth creating.

None of these alternative explanations seems compelling to me. When we attend to the tragedy of Germain missing out on greater intellectual achievements and appreciation, her life does not strike us as particularly aesthetically disvaluable and we are not directly concerned with whether her life was worth creating. Instead, when we reflect on her life, we cannot help but think that it is in one respect going well *for her* and in another respect poorly *for her*. Thus, both judgments comprising evaluative ambivalence seem to concern the prudential value of a life.

The third rival explanation accepts this conclusion but insists that the goodness of a life is determined solely by one's well-being and that the established theories of well-being can wholly account for the phenomenon of evaluative ambivalence.

Admittedly, the established theories of well-being can explain *some* forms of evaluative ambivalence. Hedonism, for example, can make sense of a divided outlook on a life which involves an abundance of joy and substantial suffering. Preference views can account for evaluative ambivalence directed towards a life in which numerous preferences are satisfied while many others are frustrated. And the objective list view could explain an apparent conflict of judgments concerning a life which is deprived of close relationships but features extraordinary intellectual achievements.<sup>6</sup>

---

<sup>4</sup> This view has been popularized by Moore (1903).

<sup>5</sup> See Parfit (1984), Blackorby, Bossert, and Donaldson (2005), and Broome (2004).

<sup>6</sup> Pummer (2017) calls such lives 'lopsided'. For concise overviews of these theories of well-being, see Gregory (2016), Heathwood (2016), and Fletcher (2016).

But all of these are instances of what we can call ‘*shallow* evaluative ambivalence’. The phenomenon that we have discovered earlier, *deep* evaluative ambivalence, is importantly different. For one thing, it seems that Germain’s life was rich in all the recognised constituents of well-being. For another, her case engenders a distinctive phenomenology. Our evaluative ambivalence has to do with judgments about the entirety of her life. We don’t feel the same way when a life is good at the beginning but bad at the end, or when it contains many pleasant experiences but also a considerable amount of suffering. Of course, calculating the net amount of well-being in such cases is not straightforward and this may be reflected in an assessment that is initially divided. But once this difficulty is behind us, there is no emotional and evaluative residue like that witnessed in Germain’s case.

The last rival explanation is subtler. It holds that the judgments comprising evaluative ambivalence track not two different kinds of prudential value, but rather two different kinds of facts about the same prudential value: well-being. In particular, one judgment tracks a *non-comparative fact* about Germain’s well-being (it was high), whereas the other tracks a *comparative fact* about her well-being (it was lower than in another possible scenario). This latter, comparative fact is what makes Germain’s life in one respect tragic. But, crucially, it does not make her life worse.

I do not find this explanation satisfactory, for two reasons. To begin with, it cannot adequately capture our patterns of intuitions in cases which have a dynamic structure, such as the following.

Suppose that you are in a café and overhear a conversation at the nearby table. “How is your son doing these days?”, one woman asks. “Same old”, the other responds, “Billy doesn’t leave the house much. He hardly ever talks to anyone and has no interest in reading—or anything, really.” The uncomfortable silence that follows is eventually interrupted. “But lately he seems to be developing some attachment to his brother. I’m curious to see where it goes.” If you are like most people, you probably think that Billy’s life is mediocre. But now suppose that the conversation progresses and you come to understand something important. Billy is not some rebellious teenager going through a difficult period, but rather a boy with a serious cognitive impairment. That newly established bond with his brother is probably as good as it could possibly be.<sup>7</sup> Once this is in view, I believe, Billy’s life doesn’t seem quite as bad for him as it initially did.

---

<sup>7</sup> This and other disability-related examples in this paper concern severe *cognitive* disabilities. For a compelling argument that negative impacts of many *physical* disabilities on one’s well-being are greatly overstated, see Barnes (2016).

Now, according to the presently considered rival explanation, our inclination to revise the initial assessment of Billy's life is a mistake. After all, the second part of the conversation between the two women has not revealed any new non-comparative facts about his well-being. But that goes against our considered judgment. In cases like this one, whether a person realizes their potential does seem to affect how good their life is for them.

The second shortcoming of this rival explanation is that it is fundamentally ill-equipped to help us make progress on the ethical puzzles addressed in the third part of this paper. As we will see, to solve these puzzles, we need something with axiological implications, whereas this rival explanation has none.

In sum, the Dual Theory appears to provide the best explanation of the phenomenon of evaluative ambivalence that we have encountered in the case of Sophie Germain and other similar scenarios.

### 3. The crib test

When philosophers seek to spell out the concept of prudential value and identify ingredients of a good life, they often appeal to a thought experiment known as *the crib test*, which works as follows.<sup>8</sup> Imagine that you are a proud parent of a newborn. As you look into the crib where the baby is peacefully sleeping, your heart fills with parental concern. You think to yourself, "I hope my baby gets to live a very good life". The kinds of things that you, a rational and benevolent caretaker, wish for in this situation, for the baby's sake, seem to be the ones which make a life worth living.

So what exactly do you wish for? Loving relationships, presumably. Achievements of various kinds. Pleasant experiences. These are the standard answers that people give and it might be tempting to stop there. But as you continue to think about that baby peacefully sleeping in the crib, I believe that you will eventually come to appreciate something important. What you care about is not just the number of relationships or achievements your child will have. You also care about the extent to which the child will realize its potential.

To put it differently, the realization of one's potential appears to be a *distinctive object of prudential concern*. These kinds of things, according to the crib test, are among the basic determinants of the goodness of a life. This is one way in which the crib test supports the Dual Theory.

But there is another lesson as well: no quantity of achievements, loving relationships, and pleasant experiences seems sufficient for an individual to have an excellent life.

---

<sup>8</sup> See Adams (1999), Darwall (2002), Feldman (2006), and Bradley (2009).

It depends, in part, on one's potential. If one's potential is high, more of these goods are needed. If it's low, fewer will suffice. In other words, the goodness of a life seems *subject-relative*: a given life, in terms of what actually happens within it, can be better for one subject than for another.<sup>9</sup>

To illustrate this feature, return to the case of Sophie Germain. Suppose that she and her sister Angelique had lives which, on the whole, contained comparable pleasant experiences, loving relationships, achievements, and disappointments, and so had the same levels of well-being. There was, however, one salient difference between them: Angelique had no extraordinary talent for mathematics. Thus, the gap between how high Sophie's well-being was and how high it could have been was larger than in Angelique's case. In virtue of this feature, Sophie's life seems overall worse than Angelique's life, even though they had equal levels of well-being.

The Dual Theory predicts this kind of subject-relativity. Because this view recognizes two basic determinants of the goodness of a life, well-being and the realization of one's potential, it implies that individuals who are tied in terms of well-being can differ in terms of the overall value of a life.

Notably, the Dual Theory also implies that it would be better for a person to have lower potential, holding fixed the way their life actually goes. Some people may initially find this corollary implausible. But I think that we should accept it upon reflection.

To begin with, let me offer a debunking explanation for the contrary intuition. There are two main mistakes that people tend to make when thinking about this issue.

The first mistake is to assume that a person could be benefited by, say, having their fabulous musical gifts removed after they are born, if they would never exercise them anyway. This would indeed be a strange implication. But the Dual Theory does not say that. On this view, it would be better for a person to be born without fabulous musical gifts, if they would never exercise them anyway. But having one's gifts removed is not like never having them in the first place. Instead, it's like being banned from playing music or suffering a career-ending accident: it does not reduce one's potential, but merely determines which of the many possible paths their life will actually take. This is how we need to think about the case of Sophie Germain as well: the sexist actions of her contemporaries did not alter what she was truly capable of,

---

<sup>9</sup> Subject-relativity should not be confused with the more widely discussed property of *agent-relativity*. A normative ethical view is agent-relative if, at least in some situations, different agents have different moral reasons for action despite having the same set of options. For canonical discussions, see Nagel (1970), Scheffler (1982), and Parfit (1984).

but merely determined which of the many possible paths her life took, and in that way affected her well-being.

But even once this is clarified, it's easy to make a second mistake: namely, to fail to hold fixed how one's life actually goes across the compared cases. There is, after all, a close connection between the potential and the actual. For instance, if someone were to lack the kind of emotional sensitivity needed to engage in loving relationships, that would typically impact not just what bonds they could have with other people, but also what bonds they would actually have. The same goes for talents. There are normally plenty of occasions for a person to put their athletic, musical, or scientific gifts to use—even if they do not end up pursuing a career in any of these fields—which would often make a positive contribution to their well-being. We may therefore implicitly associate lower potential with lower well-being even when we are explicitly asked to assume that away.

In addition to these debunking explanations, there is a positive reason to embrace the claim that it can be in one respect better for a person to have lower potential. In particular, this claim provides a natural explanation for a range of attitudes that we express towards certain kinds of news. Take the case of Michael Jordan, who retired from professional basketball in 1998 after winning a third consecutive, and sixth overall, NBA title with the Chicago Bulls. In the recent television documentary, *The Last Dance*, Jordan reflects on his spectacular career and the controversial circumstances of his decision to leave the sport. About the latter, he remarks: "It's maddening because I felt like we could have won seven".<sup>10</sup> In light of this, it seems that Jordan would be *glad* to learn that he was mistaken and in fact could have never won the seventh NBA championship, perhaps because his body could not have handled another demanding season. In a similar way, many people are *relieved* to hear that their loved one who had been missing for years has been discovered to be dead. These attitudes seem fitting and a natural explanation for why they are is precisely that having lower potential is in one respect better us. After all, a key feature which these cases share is that something that had seemed attainable—an extraordinary achievement or being reunited with a loved one—turns out not to have been in the space of possibilities after all.

With these remarks in place, let's return to the main line of the argument. The previous section provided initial support for the view that there are at least two basic prudential values. The upshot of this section—that the value of a life is subject-relative—strengthens that case. For if two lives can have the same level of well-being and yet differ in terms of how they are for the people living them, then there must be a further determinant of the life's goodness. And that determinant appears to be the degree to which an individual realizes their potential.

---

<sup>10</sup> *The Last Dance*, Episode 10.



What remains to be established is what counts as one's potential, and how well-being and the realization of one's potential combine to determine the overall goodness of a life. This is what the next two sections set out to do.

#### **4. What counts as an individual's potential?**

According to the Dual Theory, the goodness of a life is determined in part by the degree to which one realizes their potential. But what exactly counts as one's potential? There are several theoretic choice points concerning this issue.

The first thing that we should consider is whether one's potential is determined by facts about them as an individual, or rather by facts about some reference class to which one belongs, such as their species. In the latter vein, one might propose that one's potential is the maximal level of well-being that members of their species can attain.

However, this proposal is not plausible. If there was just one evaluatively significant level of potential for all humans, then many of the judgments we made earlier in our discussion would not make sense. For example, there would be nothing particularly heartening about Billy's rudimentary relationship with his brother, and no room for the claim that his life is 'fortunate' either. Likewise, in Michael Jordan's case, we would have to accept either that there was nothing tragic about him missing out on that seventh NBA title, or else that it is tragic for people like you and me that we have failed to win that many accolades. These considerations suggest that we should instead think of one's potential as determined by facts about that specific individual.

This brings us to the second choice point: which facts about the individual are relevant? Is it the level of well-being they could have easily had otherwise? Or perhaps their maximal possible well-being?

The first of these alternatives might seem attractive. While the significance of modal facts has been largely overlooked in the literature on the goodness of a life, it plays an important role in epistemology. Many authors think that the concept of knowledge includes an anti-luck condition.<sup>11</sup> Specifically, an individual fails to know that some proposition is true if they could have easily believed otherwise. Something similar could be said in the present context: the extent to which one realizes their potential is determined by the relation between their actual well-being and the level of well-being they could have easily had otherwise.

---

<sup>11</sup> See, for example, Sosa (1999), Williamson (2000), and Pritchard (2005).

However, as the case of Sophie Germain illustrates, this cannot be right. The prejudice against women, and female scientists in particular, was so entrenched in Germain's society that she couldn't have easily achieved more, received greater recognition, and avoided upsetting experiences borne of the obstacles she faced. Despite this, we intuitively judge her potential well-being to be *much* higher than her actual well-being. In light of this, and in the absence of other natural candidates, we should tentatively accept that one's maximal possible well-being is what matters.

Turn now to the third question: is the relevant concept of potential subjective or objective? Michael Jordan certainly believes that he could have won that seventh NBA title. And many of us have experienced that one romantic relationship that we had thought would last forever but which did not work out. Such instances of unfulfilled potential sting the most. It might therefore seem that our evaluations should be sensitive to one's subjective potential—in particular, the maximal possible well-being one believes to have.

Sophie Germain's case once again provides a counterexample. While Germain probably believed that her scientific achievements could have been greater, it is unlikely that she appreciated just how much larger they could have been. Through no fault of her own, of course: she is just one of the millions of women in history who have been told that they are not capable of excellence. It is preposterous to think that these are not genuine stories of wasted potential. Thus, the relevant concept of maximal possible well-being must be the objective one.

Our last question concerns the very notion of *possibility*. This question is as important as it is difficult to answer. To see that, return to the case of Billy. On one hand, because Billy is severely cognitively disabled, it seems that his life *could not have been* much better than it actually is. On the other hand, insofar as Billy could have been born without that disability, it seems that his life *could have been* much better.

To separate such cases, we may refer to the distinction between *intrinsic* and *extrinsic* potential. The former is the potential whose attainment does not require any direct or significant alteration of the individual's constitution, whereas the latter notion is more encompassing.<sup>12</sup> Thus understood, Billy's intrinsic potential is not much higher than his actual well-being, whereas his extrinsic potential is considerably higher. Of these two notions, intrinsic potential appears to be closer to what matters. After all, there *is* something heartening about Billy's life. I will assume this view for the remainder of the discussion.<sup>13</sup>

---

<sup>12</sup> For a discussion of the limitations of this distinction, see McMahan (2002) and Vallentyne (2005).

<sup>13</sup> Alternatively, we could lean into the difficulties concerning what counts as one's potential and propose that this concept is context-sensitive. On this alternative view, whether Billy's

Let's stop here for the time being. The preceding discussion suggests that one's potential is determined by facts about one's individual (rather than species), objective (rather than subjective), and intrinsic (rather than extrinsic) maximal possible well-being. There are, of course, further questions that we could ask about this concept, but the absence of a fully-fledged theory should not deter us from pressing on with our inquiry. The choice of what exactly counts as one's potential will sometimes affect the assessment of how good a particular life is, but it should not affect the substance of the claims defended in the remainder of this paper, as long we hold it fixed across the relevant comparisons.

## 5. Combining values

The Dual Theory recognizes that there are two basic prudential values: well-being and the realization of one's potential. But how exactly do they combine to determine the overall value of a life? I would like to restrict the present discussion to two simple models which represent two different ways of conceptualising the significance of wasted potential.

According to *the Addition Model*, to determine the overall value of a life, we need to add the value of realized potential to the contribution made by well-being, where the former is represented by a non-negative number. This is a 'glass half-full' view: one gets bonus points for realizing one's potential, and more points for realizing it more fully.

By contrast, *the Subtraction Model* determines the overall value of a life by subtracting the disvalue of unrealized potential from the value of well-being, where the former is represented by a non-negative number. This is a 'glass half-empty' view: one loses points for not realizing one's potential fully, and that penalty is greater the bigger the potential gap.

Considering that addition and subtraction are inverse mathematical operations, these two views will have the same implications in a wide range of cases. However, in virtue of how they conceptualise the significance of realizing one's potential, these models

---

intrinsic or extrinsic potential is relevant depends on the context in which we are assessing the goodness of his life. This would be an instance of a broader phenomenon of the context-sensitivity of counterfactuals. (See Lewis 1986 for a canonical discussion.). This alternative view is not pursued here any further, but it is worth noting one of its implications. If what counts as one's potential were context-sensitive, and the Dual Theory were true, then the overall goodness of a life would also be context-sensitive—and that would have significant upshots for normative ethics.

can diverge in their assessment of lives with barely positive levels of well-being and lives with barely negative levels of well-being.

To see that, first consider a life which contains no positive ingredients of well-being (pleasures, loving relationships, achievements, etc.) and a small amount of negative ingredients of well-being (intermittent pain, for example). The Addition Model allows for such a life to be overall worth living, as long as one's maximal possible well-being is not much higher, and so one's potential is realized to a large extent. This strikes me as deeply counterintuitive. By contrast, according to the Subtraction Model, a life with a negative well-being level is never overall worth living.

On the other hand, the Subtraction Model implies that unless we make *enough* of our talents and opportunities, our lives will fall short of being overall worth living. Thus, a life with a positive level of well-being could be not worth living overall. I find this reverse implication more plausible.

To begin with, consider a human living the life of a domesticated mouse which, let's assume, is short, monotonous, and deprived of sophisticated pleasures, achievements, and deep relationships. Intuitively, although this life would have a positive level of well-being, it does not seem worth living for a typical human whose potential is high.<sup>14</sup>

Moreover, unless we accept something like the Subtraction Model, we cannot make sense of our evaluative attitudes towards smaller-scale events, such as the following.

*Marathon:* Eve is capable of qualifying for the Olympic Games in the marathon but, instead of attending the Olympic Trials, she participates in a local race and wins it by being only slightly faster than the amateur second-place finisher.

Whatever contribution this accomplishment makes to Eve's well-being, it is overshadowed by the badness of failing to realize her athletic potential in that part of her life. Crucially, that is not to say that it would have been better for Eve not to participate in that amateur race at all. After all, the gap between no achievement and a great achievement is larger than the gap between a small achievement and a great achievement. But such barely good performances are not enough to render a life with extraordinary potential overall worth living. These considerations lead me to accept the Subtraction Model.

---

<sup>14</sup> McMahan (1996, p. 9) expresses a similar sentiment in the process of arguing that 'fortune', not well-being, should be the focus of principles of justice: "A dog, for example, has a relatively low level of well-being. But while a normal human adult with a comparable level of well-being would be very badly off, the dog may well be flourishing".

Admittedly, some people do not share my intuitions about these cases. They find too extreme the idea that a life with a positive level of well-being can be overall not worth living, even if it realizes its potential to a very low extent. To avoid this implication, they could opt for an asymptotic version of the Subtraction Model, on which unrealized potential always makes a life go worse but it never makes it overall not worth living.<sup>15</sup> This alternative view would be in one way stronger: it would be easier to accept for most people. In another way, it would be weaker: some of its explanatory power would be lost. In particular, the Dual Theory would no longer deliver a solution to the puzzle discussed in the next section, concerning standards for a life worth living in animals. When weighing these theoretical virtues, I find myself partial to greater explanatory power and thus prefer the simple Subtraction Model. But the reader can choose otherwise. All other claims would still follow.

## 6. Humans, animals, and lives worth living

The preceding discussion provided initial motivation for the Dual Theory and examined some of the ways in which it can be spelled out. What I want to do in the remainder of the paper is to put the Dual Theory to test and show that it illuminates three ethical puzzles.

In this section, I argue that the orthodox view which holds that well-being is the sole determinant of the goodness of a life faces a *trilemma* regarding standards for a life worth living for humans and non-human animals.<sup>16</sup> We can avoid this problem, however, if we accept the Dual Theory.

The trilemma arises when those who accept the orthodox view are confronted with the following, simple question: What counts as a life worth living for humans and animals?

The first option available to them is to postulate that there is a single threshold for a life worth living for humans and animals, and to set this threshold at a ‘relatively low’ level of well-being. This view captures the intuitive thought that at least some animals have good lives: your friend’s Labrador, perhaps, or the cows grazing peacefully in the Swiss Alps. Notice, however, that this view also implies that a cow-like life—short, monotonous, and deprived of sophisticated pleasures, significant

---

<sup>15</sup> For example: Value of a Life = Well-being<sup>2</sup>/Potential.

<sup>16</sup> A background assumption needed to generate this trilemma is *welfare invariabilism*, which is the view that the same theory of well-being is true of every subject capable of well-being. That is, the same list of basic goods (such as pleasures, achievements, friendships) applies to all humans and non-human animals who are capable of welfare, even if some of these goods are in fact inaccessible to them. For a compelling defence of this view, see Lin (2018).

achievements, and deep relationships—would be worth living for a typical human whose potential is high. As noted earlier, this seems difficult to accept.<sup>17</sup>

The second option is to uphold the commitment to a single threshold but to set it at a ‘relatively high’ level of well-being instead. This avoids the implication that a cow-like life would be worth living for a typical human. However, this obviously comes at the expense of capturing the first intuition that some animals have good lives. This may not seem like a high cost at first. But if we also grant the claim that there is a moral reason to prevent individuals who would have lives not worth living from coming into existence, then it follows that we have a moral reason to bring about the extinction of many, if not all, animal species.<sup>18</sup> This, too, is deeply counterintuitive.

The remaining alternative is to propose that there are two standards for a life worth living: one for humans, at a ‘relatively high’ level of well-being, and another for animals, at a ‘relatively low’ level of well-being.

But this option is also unsatisfactory, for at least a couple of reasons. For one thing, it is speciesist: it discriminates among individuals on the basis of their species membership *alone*.<sup>19</sup> Two individuals which have the same level of well-being should not differ with respect to whether they have lives worth living just because they belong to different species.

For another, this last answer has some unpalatable implications for distributive ethics, at least if we think that animals are within the scope of egalitarian concern. For the sake of precision, let’s say that well-being can be represented on a ratio scale, and that the threshold for a life worth living is at 0 units of well-being for humans, and at -10 units of well-being for animals. Suppose that there is a human, Ada, whose well-being is at -1 and that there is a dog, Fido, whose well-being is at -9. In this case, we would have an egalitarian reason to redistribute well-being from Ada to Fido. But this means that we would have a reason to redistribute well-being from an individual who has a life not worth living to an individual who has a life worth living. That’s not plausible.

---

<sup>17</sup> Two clarifications. First, a cow-like life need not literally be the life of a cow. It could be a life that is equivalent to that of a cow in terms of the level of well-being. Second, such a life does not obviously seem not worth living for an atypical human whose potential is low (perhaps in virtue of a severe cognitive disability.) The argument presupposes only that it is not worth living for a typical human whose potential is high.

<sup>18</sup> This claim is one part of the so-called *Procreation Asymmetry*. See McMahan (1981) and Frick (2020).

<sup>19</sup> This term was popularized by Singer (1975).

We can avoid this trilemma if we recognize that equating the goodness of a life with the level of well-being is a mistake. A cow might have the same level of well-being as a typical human living a cow-like life, but these two individuals are likely to differ in terms of the degree to which they realize their potential. The cow's life could not have been much better than it is, whereas the human's life could have been. This allows us to explain how, despite having the same levels of well-being, a cow could have a life worth living and a human could have a life not worth living. Thus, the Dual Theory delivers a compelling response to our trilemma. It allows us to have a single threshold for a life worth living for all beings. At the same time, it avoids implying that no animals have lives worth living or that a cow-like life would be worth living for a typical human.

In closing, let's briefly address two issues. First, why think that animals typically realize their potential to a higher degree than humans? Humans are generally equipped with greater and more numerous capacities than non-human animals, and therefore can access a wider range of goods. For instance, humans are capable of valuing their bonds with others and experiencing a broader array of emotions regarding these bonds, which allows them to participate in highly intimate relationships. Likewise, humans seem to be uniquely equipped with the awareness of what is worth pursuing and the grit to succeed in doing so, which are essential ingredients for the greatest achievements.

Thanks to these capacities, most humans can have truly extraordinary lives. We can get an idea of what is possible by reflecting on our life's best moments and imagining that our entire lives were like that. Our maximal possible well-being is indeed very high. But, as the experiments in living of billions of our ancestors and contemporaries teach us, it is extremely difficult to flourish in all respects. Thus, human lives which brush against their maximal possible well-being are exceedingly rare. By contrast, for non-human animals, whose capacities are much more limited, the task of reaching one's potential is easier. We should therefore expect animals to realize their potential to a greater degree than humans.

Second, isn't there a rival non-orthodox solution to the trilemma identified above? On the Dual Theory, the realization of one's potential and well-being jointly determine the goodness of a life. But one could propose a more radical view, on which the realization of one's potential is the sole determinant of the life's value. Both of these views avoid the trilemma identified above.

However, the radical view does so at a high price. First, in virtue of its monistic character, the radical view cannot account for the phenomenon of evaluative ambivalence identified earlier. Second, this view judges animal lives to be implausibly good: for instance, it implies that the life of a fortunate dog (low well-being, low potential) is just as good as the life of a fortunate human (high well-being, high potential).

potential). Third, the radical view has some unpalatable implications for distributive matters: for example, it implies that we should benefit Aisha rather than Bertrand if Aisha has a much higher level of well-being than Bertrand, but realizes her potential only to a slightly lesser degree than him. The Dual Theory does not run into these problems.

## 7. The significance of a life's shape

The Dual Theory also helps us resolve a puzzle concerning the following question: What is the relation between how well one's life is going at any particular time and how well it goes overall?

The following, simple answer might initially seem attractive. To determine how well a life goes overall, we can simply assess how well it is going in each period and then put these values together. Moreover, the temporal features of a life, such as the order of events, do not affect the prudential value of a life. Of course, a vigorous game of tennis followed by an indulgent lunch makes for a more pleasant day than the converse. But, the thought goes, it makes no difference which days are happy and which are miserable, as long as their number is the same.

However, this simple view is challenged by the following pair of cases.

*Improvement.* Your life begins in the depths and takes an upward trend: misery in the early years, mixed fortune in midlife, followed by flourishing in old age.

*Deterioration.* Your life begins at the heights and takes a downward trend: an early period of flourishing, mixed fortune in midlife, followed by misery in old age.

These lives are permutations of each other—in the sense that, for every temporal segment with some level of well-being in one life, there is exactly one such segment in the other life, and so on—so the simple view regards them as equally good. However, many people have a strong intuition that the improving life is better than the deteriorating life.<sup>20</sup> The challenge is to find a compelling theoretical rationale for this judgment.

Although this puzzle has generated a substantial literature, it will be instructive to start with my positive proposal and only then turn to the shortcomings of the existing

---

<sup>20</sup> This intuition is shared by Slote (1983), Bigelow et al. (1990), Velleman (1991), Kamm (2003), Portmore (2007), Temkin (2012), Kauppinen (2012), Glasgow (2013), Dorsey (2015), and Hirose (2015).



views. In essence, I argue that an improving life typically features a greater degree of potential-realization than a deteriorating life, and so typically it is overall better even if the total amounts of well-being are the same. This explains why we are inclined to judge that the improving life is better in the original, abstract case.

To illustrate this, consider two ways in which a philosophical career could turn out.

*Early Breakthrough:* Adam writes his most original and insightful article shortly after finishing his graduate studies. His next publications continue to make valuable contributions to philosophical debates, but their quality gradually diminishes.

*Late Breakthrough:* Adam begins his career with a run-of-the-mill article, but the originality and insight of his work steadily improve.

These two lives, let's assume, are otherwise identical.<sup>21</sup> The second life seems better than the first. Why?

Take Early Breakthrough. There is just something tragic about a situation in which the greatest achievements take place early on in Adam's life and are not replicated further down the road. A natural explanation of this tragedy draws on the same kinds of considerations that we have been discussing throughout: the significance of realizing one's potential. The fact that Adam writes a terrific essay early on in his life reveals that he is already capable of producing great works of philosophy at that stage. His *momentary potential*, as it were, is already high. When his next paper is a lesser achievement, this amounts to more than just an over-time deterioration in terms of *momentary well-being*. In addition, the gap between what Adam could accomplish and what he actually accomplishes grows, and with it grows the gap between his well-being and his potential.

Consider Late Breakthrough now. The fact that Adam's career starts with a minor paper suggests that he does not yet have what it takes to write a brilliant article. Consequently, the gap between what he achieves and what he could achieve at that time is small. Moreover, it is natural to assume that the subsequent improvement in the quality of Adam's work mirrors the increase in his momentary potential.

Thus, although these two lives have the same levels of *lifetime well-being*, they differ in terms of the level of *lifetime realization of one's potential*. It is greater in Late Breakthrough

---

<sup>21</sup> Of course, this involves a high degree of idealization. Discovering that your work is getting worse would be a source of distress for many people. But we need to assume that away or suppose that Adam just does not give any thought to this matter.

than in Early Breakthrough. This explains why the former life seems better than the latter.

An analogous gloss can be given for a comparison between two lives which differ in terms of when one's most prized relationships occur.

*Lost Love:* In his youth, Alex has a number of highly fulfilling relationships. Over the years, however, these attachments crumble and are replaced by other, more superficial connections.

*Found Love:* Alex's youthful relationships are relatively superficial. Over time, however, he makes new, highly fulfilling connections.

In *Lost Love*, the fact that Alex has highly fulfilling relationships early in his life reveals that he already has the kinds of emotional sensitivity, sense of respect, and the ability to trust another that are needed to engage in and sustain the most valuable relationships. When his next attachments are of lesser quality, the gap between how high his well-being is and how high it could be, at any given time, starts to grow. By contrast, in *Found Love*, the relationships that Alex actually has at any given time presumably correspond to the relationships that he could have at that time. Thus, these two lives differ in terms of the level of lifetime realization of one's potential. This explains why *Lost Love* strikes us as worse than *Found Love*.

Crucially, this account does not imply that a life which improves in terms of well-being is *always* better than a life which features a corresponding deterioration. There are at least two kinds of scenarios in which it isn't.

First, note that not all cases of early breakthrough strike us as tragic. Mathematicians, for example, tend to make the most substantial contributions to their field at a relatively young age. Indeed, the highest honour a mathematician can receive, the Fields Medal, has an age limit (the recipient must be under forty), but this is hardly a constraint in selecting the winner. Similarly, most athletic careers end in one's thirties. These life stories do not seem as upsetting as that of a philosopher who fails to make a substantial scholarly contribution in the second half of their life.

Why? Evidently because of the difference in terms of what the subjects are capable of at any given point in their life. Since the abilities of mathematicians and athletes typically peak in their late twenties or early thirties and then deteriorate, they do not fail to realize their potential if they fail to replicate the achievements of their youth later in life. Arguably, philosophers are not like that: insofar as there is a peak, it typically comes later in life, and thus a story of an early breakthrough is bound to be more upsetting.

Second, there are cases in which the contrasting patterns of well-being are attributable to changes in one's pleasant experiences rather than achievements or relationships. Dale Dorsey (2015, p. 139) considers one such scenario.

[C]ompare two different experiments. The first person is hooked up to a pleasure-stimulating computer with software designed to start his life at a neutral level: no pleasure, no pain. Gradually, say, twice a year, the pleasure is increased in a linear fashion. The second is precisely the opposite: the person's machine starts out with quite a lot of pleasure. Gradually, also twice per year, the software decreases the pleasure in a linear fashion, such that both people, over the course of their lives, will generate the same amount of sensory pleasure.

Intuitively, neither of these two lives seems better than the other. My account accommodates this verdict. In the situation imagined by Dorsey, there is no reason to suppose that there is a difference in terms of the subjects' realization of their potential. Thus, the Dual Theory judges these lives to be equally good.

To summarise, here is how we should think about the puzzle at hand. The shape of a life indicates (but does not guarantee) the presence of something else that affects the goodness of a life: whether one's potential has been realized to a high extent. It is in virtue of differences in this respect that an improving life is *typically* better than a deteriorating life. This explains why we are inclined to judge an improving life to be better than a deteriorating life in the original case. Call this *the Potential Account*.

This solution avoids the shortcomings of three popular accounts offered in the literature. According to Michael Slote's (1983) *Temporal Location Account*, an upward-sloping life is better than a downward-sloping life because what happens later in life has greater weight in determining the life's overall value. That is, any given good (say, an achievement or a pleasant experience) contributes more to the goodness of a life if it is present in one's adulthood or old age rather than in one's youth.

*The Pattern Account* offers a different take. It regards improvements in terms of temporal well-being as in themselves good for a person and corresponding deteriorations as in themselves bad for them. Frances Kamm (1998; 2003), Larry Temkin (2012), Joshua Glasgow (2013), and Iwao Hirose (2015) are all in this camp.

Finally, according to *the Narrative Account*, an upward-sloping life is better than a downward-sloping life when and because it features certain narrative relations between events in one's life, such as when one suffers a misfortune and then learns from it, or when an effort is turned into a success rather than a waste. David Velleman (1991) has defended this view, and its variants have been endorsed by Douglas Portmore (2007), Antti Kauppinen (2012), and Dale Dorsey (2015).

An undeniable attraction of the first two views is their simplicity. But it is also their greatest flaw: they *prove too much*. In particular, they predict that an improving life is always better than a deteriorating life. But, as we have seen, this is not the right verdict.

The Narrative Account is closer to the truth in that it recognizes that an improving life is only typically better than a deteriorating life. But Velleman's stated account *proves too little*. It cannot capture the relevant intuitions in some central cases, such as the comparison between Found and Lost Love, and between Late and Early Breakthrough. In these scenarios, one of the compared lives seems better than the other even if we stipulate that they do not differ in terms of whether the agent has learned from their past mistakes or turned their past efforts into an eventual success. Moreover, it is far from obvious how the Narrative Account could be extended to cover such cases. Velleman does not offer a unified account of the significance of the relations between discrete events that he discusses. This is in stark contrast to the Potential Account which explains our intuitions about the relevant cases in a theoretically unified way. On this view, an improving life is better than a deteriorating life when and because these alternatives differ in terms of the degree to which one's potential is realized.

## 8. Badness of death

The Dual Theory can also illuminate a long-standing ethical puzzle concerning the badness of death. Many people share the belief that death is typically bad for its subject, and one of the things we want to know is why and to what extent this is so.

Nagel (1979) and others who have addressed this question quickly rejected the implausible views that death is bad for us simply because being dead or dying are bad for us. This left them with the idea that death is bad because it *deprives* us of future goods. One popular version of this view is known as *the Life Comparative Account*.<sup>22</sup> It says that to determine the extent to which death is bad for someone who dies in an accident, we should compare the amount of well-being they actually had with the amount of well-being they would have had, were they not to die in the accident.

One thing that we can notice straightaway is that the Dual Theory dovetails with the Life Comparative Account. For instance, both views capture the profound sense of tragedy that accompanies the death of a teenager who 'had a whole life ahead of him'. And both views recognise that, in some cases, an earlier death could be a blessing.

---

<sup>22</sup> This label is due to McMahan (2002, p. 105). See also Feldman (1991), Broome (1993), and Bradley (2009).

However, the relation between these two accounts is deeper than that: one explains the other. According to the Dual Theory, the badness of death is just one instance of a more general phenomenon, the badness of unrealized potential. To put it in a slogan form, being run over by a bus is like failing to catch a bus. It makes us miss out on various opportunities presented to us by life. Of course, the tragedy of a premature death is typically vastly greater than the setback caused by failing to catch a bus. But that's because the former makes us miss out on *all* opportunities, not just some, and thus leaves a greater chunk of our potential unrealized.

## 9. Concluding remarks

I have argued that how good a life is for someone is determined jointly by their level of well-being and the degree to which they realize their potential, where the latter is to be understood in terms of how close one is to their maximal possible well-being. In particular, I have claimed that this Dual Theory provides the best explanation of the phenomena of evaluative ambivalence and subject-relativity, and that it illuminates ethical puzzles concerning standards for a life worth living for animals, the significance of a life's shape, and the badness of death.

Of course, there is more work to be done if the Dual Theory is to supplant the orthodox view on which the goodness of a life depends exclusively on what actually happened within it. The arguments of this paper notwithstanding, the Dual Theory remains an extraordinary claim, and such claims require extraordinary evidence. Further analysis may also show that some of the details of this view—for example, concerning what counts as one's potential and how to combine well-being and the realization of one's potential—should be filled in alternative ways. But I hope that my discussion has at least provided a blueprint for this larger project and motivation for pursuing it.

In closing, I would like to address two worries which are often raised about this account. They both concern how we should direct our lives, given what know and don't know about our potential.

The first issue is this. According to the Dual Theory, prudence demands that we make the most of our talents and capabilities. Indeed, as we saw in Section 5, unless we make enough of our talents and capabilities, our lives will fall short of being worth living. But we also know that, for example, athletes and musicians tend to make tremendous sacrifices in their lives: often, if you are ten years old and not already practicing something, you have little chance of world-class success. Moreover, if you do not devote dozens of hours every week to your discipline, you are unlikely to get very far. To make matters worse, only a tiny proportion of those who put in the work will make it to the main stage. Now, we certainly cheer on people who choose to take

this path, but a theory of prudence which condemns people to it may seem objectionably demanding.

However, the upshot of the Dual Theory is more nuanced. While my discussion has centred on athletic and intellectual achievements, it is important to keep in mind that any plausible theory of well-being recognises a range of other prudential goods: relationships, pleasures, and knowledge, among others. In virtue of this, we have to be sensitive not only to whether an individual realizes their potential for achievements, but also to whether they realize their potential for these other goods. At one extreme, if running or doing philosophy are truly the only things that make you happy and you show great promise in these disciplines, prudence may well demand that you go all-in. At the other end, if you are just above average and these pursuits would eat up all of your time, then it may be better to choose a safer career and live a more balanced life.

One might also worry that as we expand our understanding of the goodness of a life, it becomes less transparent what is in our interest at any given time and how our lives are going as a whole. This is because now we need to know not just facts about the actual world, but also about various possible worlds. And this, in turn, might be taken to have a paralysing effect on the ability of prudence to inform our lives.

This concern is overstated. Even the orthodox view, on which the goodness of a life is determined solely by the level of well-being, gives us at most rough guidance regarding matters of self-interest. For example, we would all be hard-pressed to report our lifetime hedonic score as of this morning. Likewise, the primary ambition of the Dual Theory is not to guide us through every single decision that we have to make as we go through life. Instead, its aim is to help us better understand what makes life good. I believe it is successful in this regard.

## References

- Adams, R. M. (1999). *Finite and Infinite Goods: A Framework for Ethics*. New York: Oxford University Press.
- Barnes, E. (2016). *The Minority Body: A Theory of Disability*. Oxford: Oxford University Press.
- Bigelow, J., Campbell, J., & Pargetter, R. (1990). Death and Well-being. *Pacific Philosophical Quarterly*, 71(4): 119-140.
- Blackorby, C., Bossert, W., and Donaldson, D. (2005). *Population Issues in Social-Choice Theory, Welfare Economics and Ethics*. New York: Cambridge University Press.
- Bradley, B. (2009). *Well-being and Death*. Oxford: Oxford University Press.

- Broome, J. (2004). *Weighing Lives*. Oxford: Oxford University Press.
- Dorsey, D. (2015). The Significance of a Life's Shape. *Ethics*, 125(2): 303-330.
- Feldman, F. (1991). Some Puzzles About the Evil of Death. *The Philosophical Review*, 50(2): 205-227.
- Feldman, F. (2006). *Pleasure and the Good Life: On the Nature, Varieties, and Plausibility of Hedonism*. New York: Oxford University Press.
- Fletcher, G. (2016). Objective list theories. In G. Fletcher. *The Routledge Handbook of Philosophy of Well-being*. New York: Routledge.
- Frick, J. (2020). Conditional Reasons and the Procreation Asymmetry. *Philosophical Perspectives: Ethics*, 34(1): 53-87.
- Frick, J. & Lerner, A. (ms). Speciesism and Acceptable Lives.
- Glasgow, J. (2013). The Shape of a Life and the Value of Loss and Gain. *Philosophical Studies*, 162(3): 665-682.
- Gregory, A. (2016). Hedonism. In G. Fletcher. *The Routledge Handbook of Philosophy of Well-being*. New York: Routledge.
- Harman, E. (2009). "I'll be glad I did it" reasoning and the significance of future desires. *Philosophical Perspectives*, 23: 177-99.
- Heathwood, C. (2016). Desire-fulfillment theory. In G. Fletcher. *The Routledge Handbook of Philosophy of Well-being*. New York: Routledge.
- Hirose, I. (2015). Intra-personal aggregation. In his *Moral Aggregation*. Oxford: Oxford University Press.
- Kamm, F. (1998). *Morality, Mortality Vol. 1: Death and Whom to Save From It*. New York: Oxford University Press.
- Kamm, F. (2003). Rescuing Ivan Ilyich: How we live and how we die. *Ethics*, 113(2): 202-233.
- Kauppinen, A. (2012). Meaningfulness and Time. *Philosophy and Phenomenological Research*, 84(2): 345-377.
- Lewis, D. (1986). *On the Plurality of Worlds*. Oxford: Blackwell Publishers.
- Lin, E. (2018). Welfare Invariabilism. *Ethics*, 128(2): 320-345.
- McMahan, J. (1981). Problems of Population Theory. *Ethics*, 92(1): 66-127.
- McMahan, J. (1996). Cognitive Disability, Misfortune, and Justice. *Philosophy and Public Affairs*, 25(1): 3-35.
- McMahan, J. (2002). *The Ethics of Killing*. New York: Oxford University Press.

- Moore, G. E. (1903). *Principia Ethica*. Cambridge: Cambridge University Press.
- Nagel, T. (1970). *The Possibility of Altruism*. Princeton: Princeton University Press.
- Nagel, T. (1979). Death. In his *Mortal Questions*. Cambridge: Cambridge University Press.
- Parfit, D. (1984). *Reasons and Persons*. Oxford: Oxford University Press.
- Portmore, D. (2007). Welfare, achievement, and self-sacrifice. *Journal of Ethics and Social Philosophy*, 2(2): 1-28.
- Pritchard, D. (2005). *Epistemic Luck*. New York: Oxford University Press.
- Pummer, T. (2017). Lopsided Lives. In Mark Timmons (ed.), *Oxford Studies in Normative Ethics*, vol. 7. Oxford: Oxford University Press.
- Scheffler, S. (1982). *The Rejection of Consequentialism*. Oxford: Clarendon Press.
- Setiya, K. (2016). Retrospection. *Philosophers' Imprint*, 16(15): 1-15.
- Singer, P. (1975). *Animal Liberation: A New Ethics for Our Treatment of Animals*. HarperCollins.
- Slote, M. (1983). Goods and Lives. In his *Goods and Virtue*. Oxford: Clarendon Press.
- Smuts, A. (2017). *Welfare, Meaning, and Worth*. London and New York: Routledge.
- Sosa, E. (1999). How Must Knowledge Be Modally Related to What Is Known? *Philosophical Topics*, 26(1/2): 373-384.
- Temkin, L. (2012). *Rethinking the Good*. New York: Oxford University Press.
- Vallentyne, P. (2005). Of Mice and Men: Equality and Animals. *The Journal of Ethics*, 9(3/4): 403-433.
- Velleman, D. (1991). Well-being and Time. *Pacific Philosophical Quarterly*, 72(1): 48-77.
- Wallace, R. Jay. (2013). *The View from Here: On Affirmation, Attachment, and the Limits of Regret*. New York: Oxford University Press.
- Williamson, T. (2000). *Knowledge and its Limits*. Oxford: Oxford University Press.
- Wolf, S. (2010). *Meaning in Life and Why it Matters*. Princeton: Princeton University Press.